PRD — CaseFlow MCP

MCP-first unified legal workflow platform for attorneys

---

1) Document Information

Product Name: CaseFlow MCP
Version: v1.0 PRD
Date: June 24, 2026
Prepared For: Open-source MVP / early product build
Product Type: Legal workflow intelligence platform with MCP connectors, AI-assisted intake, timeline generation, and legal knowledge workflows

---

2) Executive Summary

Lawyers and legal teams work across fragmented systems such as Outlook, Gmail, Teams, Slack, SharePoint, Google Drive, Word, and matter-management systems, causing requests, documents, deadlines, and important facts to become scattered across channels. This fragmentation slows triage, increases operational stress, reduces visibility, and contributes to the risk of missed deadlines and poor matter coordination.

At the same time, the legal market is increasingly adopting AI tools that integrate into existing workflows rather than forcing lawyers to switch systems. Vendors such as Wordsmith emphasize legal work embedded inside Slack, Teams, Outlook, Gmail, Word, and connected repositories, while Filevine positions its stack as a broader legal operating system spanning intake, workflows, documents, timelines, deadlines, and AI-assisted execution.

CaseFlow MCP is an **MCP-first legal operations layer** that connects common communication and document systems, converts unstructured legal communications into structured work, extracts tasks/events/deadlines, builds matter timelines, and provides grounded legal assistance with approval-first workflow controls. MCP is specifically relevant because it is an open protocol designed to connect LLM applications to external resources, tools, and workflows in a standardized way.

---

3) Problem Statement

Legal professionals face recurring operational problems that are not solved by generic AI chatbots:

3.1 Fragmented communication

Legal requests and updates arrive across inboxes, chat systems, shared drives, and ad hoc channels, which forces lawyers to manually reconstruct context and triage work. Checkbox explicitly highlights legal intake fragmentation across email, Slack, Teams, Jira, and forms as a major operational problem, while legal workflow case studies also show that critical documents and updates often get buried in scattered threads.

3.2 Manual intake and triage

Legal teams receive a large volume of unstructured requests that require classification, routing, prioritization, and ownership assignment. Vendors in this category increasingly focus on legal intake automation and triage because rising work demand and limited legal headcount make manual triage inefficient and unsustainable.

3.3 Missed deadlines and calendaring risk

Missed legal deadlines remain a major malpractice and operational risk, and manual tracking across emails, calendars, and spreadsheets creates failure points. Legal calendaring and legal operations commentary repeatedly describe deadline management as a severe risk area for litigators and legal teams.

3.4 Timeline reconstruction is expensive and error-prone

Legal work often requires manually rebuilding a chronology of events from communications, filings, attachments, and notes. Filevine's product set includes chronology/timeline-driven intelligence, showing that turning unstructured case information into a structured timeline is a real and valuable legal workflow.

3.5 Knowledge is trapped in repositories

Legal teams repeatedly search across prior matters, templates, policies, document repositories, and knowledge bases to answer recurring questions. Wordsmith explicitly markets the value of connecting repositories like SharePoint and Google Drive into a single, searchable legal brain available from the tools users already use.

3.6 AI adoption in legal requires trust and control

The ABA has made clear that lawyers using generative AI must still satisfy obligations related to competence, confidentiality, communication, supervision, and fees. This means legal AI products require human oversight, careful handling of client data, and transparency into outputs and source grounding.

---

4) Product Vision

CaseFlow MCP will be the unified legal workflow layer that turns fragmented communications into structured legal work across the tools attorneys already use.

The platform will:

- connect communication and document systems through MCP and standard APIs,
- triage and structure inbound legal requests,
- extract events, deadlines, and tasks from threads and attachments,
- maintain a matter-centric legal memory,
- generate evidence-backed timelines,
- provide grounded answers and drafting support using approved knowledge sources,
- and execute actions under approval-first and permission-aware controls.

---

5) Product Goals

Primary Goals

1. **Reduce legal communication overload** by centralizing email/chat-derived legal work into one command center.
2. **Improve intake and triage speed** through automated classification, routing, and summarization.
3. **Reduce deadline risk** by extracting trigger events and surfacing calendar-relevant obligations.
4. **Create a matter timeline automatically** from communications and documents.
5. **Enable trustworthy legal AI workflows** through source-grounded retrieval, role-aware access, and human approval.

Secondary Goals

1. Build an extensible **open-source MCP-based legal workflow platform**.

2. Support future integrations with matter systems and CRMs via MCP/API.
3. Create a differentiated portfolio project demonstrating connectors, RAG, extraction pipelines, workflow orchestration, and legal-specific governance.

---

6) Non-Goals (v1)

The v1 platform will **not** attempt to replace a full legal practice-management suite. That means the following are explicitly out of scope for the initial release:

- Trust accounting and full legal accounting systems, which are specialized and heavily operational.
- Full end-to-end billing/invoicing engines, although event capture may later support time-entry suggestions.
- Full court-rule computation for every jurisdiction, because court calendaring logic is broad and jurisdiction-specific.
- Full document redlining and enterprise-grade clause negotiation across all contract types. Wordsmith and Filevine show this is a broad domain that can be built later, but it should not be the v1 wedge.
- Consumer WhatsApp scraping or unsupported communication ingestion methods. Production-grade integrations should rely on supported APIs and enterprise connectors. This aligns with the overall integration-first and governance-oriented direction of the product.

---

7) Target Users

Primary User Segments

7.1 Small and mid-size law firms

These firms often suffer from tool fragmentation, heavy communication load, and limited admin support, making them strong candidates for communication intelligence and workflow automation. Legal industry reporting notes pressure from growing caseloads, financial strain, and the need for simplified, integrated processes.

7.2 In-house legal teams

In-house legal teams increasingly receive requests from business users through Slack, Teams, email, and forms, making legal intake and triage a core problem space. Checkbox and Wordsmith both directly address embedded intake and conversational legal workflows in workplace tools.

7.3 Litigation-heavy or investigation-heavy practices

Practices with high needs for timelines, evidence chronology, and deadline sensitivity will benefit most from event extraction and timeline assembly. Filevine's emphasis on chronology and deposition intelligence strongly supports this use case.

---

8) Personas

Persona 1: Litigation Associate

Needs to monitor multiple matters, reconstruct fact patterns quickly, and avoid missing filing dates or response triggers. Manual calendaring and timeline building create heavy cognitive overhead.

Persona 2: In-House Counsel

Receives legal requests from different business teams through Slack, email, and Teams, and wants a way to automatically structure requests, route routine questions, and retrieve policy-backed answers.

Persona 3: Paralegal / Legal Operations Manager

Needs operational visibility into requests, deadlines, document links, assignments, and matter activity, with reduced dependency on manual tracking across disconnected systems.

Persona 4: Legal Tech / Innovation Lead

Wants AI capability but must ensure privacy, access controls, auditability, and human oversight, in line with legal ethics expectations.

---

9) User Jobs to Be Done

1. "When a legal request arrives in email, Teams, or Slack, help me **identify what it is, how urgent it is, and who should own it**."

2. “When I open a long thread, show me the **summary, key facts, deadlines, tasks, and legal risk signals**.”
3. “When I’m working a matter, help me **see a reliable, source-linked timeline of events**
4. “When I ask a legal question, answer using **approved internal knowledge and linked sources**, not generic internet content
5. “When the system finds a deadline or action item, make it easy to **create a task, assign it, and track it**.”

---

10) Product Scope (v1)

The initial v1 scope will focus on the highest-value wedge:

v1 Wedge:

Communication Intelligence + Intake/Triage + Timeline + Deadline Extraction

This wedge is supported by strong market demand around embedded legal workflows, intake triage, legal communication overload, and timeline/deadline risk.

---

11) Core Features

11.1 Connector Hub (MCP-first integration layer)

Description

A connector layer that ingests data and enables actions across communication and document systems using MCP where possible and standard APIs where needed. MCP is an open protocol intended to expose resources, tools, and workflows to AI applications in a standardized way, and Microsoft also now provides an enterprise MCP server model for Microsoft Graph scenarios.

v1 Connectors

- Outlook / Microsoft 365 via Graph-compatible patterns.
- Gmail (API-based connector). Wordsmith’s inbox workflows reflect strong demand for email-native legal workflows across both Outlook and Gmail.
- Teams for legal intake and collaboration workflows.

- Slack for intake, legal Q&A, and message triage.
- SharePoint / Google Drive for approved knowledge retrieval.

Out of scope for initial connector set

- Unsupported consumer messaging integrations.
- Full matter-system write-back connectors unless kept lightweight.

These exclusions keep the MVP realistic while aligning with the integration capabilities already proven valuable in legal AI products.

11.2 Legal Intake & Triage Engine

Description

Convert unstructured inbound communication into structured legal work items.

Inputs

- emails,
- Teams messages,
- Slack messages,
- forwarded files,
- intake form submissions.

Outputs

- matter/client association,
- request category,
- urgency,
- owner recommendation,
- summary,
- next recommended action.

Example categories

- contract review,

- litigation support,
- privacy/data inquiry,
- employment issue,
- commercial dispute,
- compliance question,
- policy question.

This reflects how legal intake platforms route requests by work type, specialization, and operational importance.

11.3 Communication Summarization Layer

Description

Summarize threads and conversations inside the legal context of the matter.

Outputs

- concise summary,
- key facts,
- open questions,
- legal issues signaled,
- action items,
- unresolved requests.

Requirements

- summary must include source links to the originating thread/message/document,
- summary must distinguish between confirmed facts and inferred issues,
- summary must identify unknowns or missing documents.

These controls align with legal trust expectations and source grounding requirements implied by legal AI ethics concerns.

11.4 Event, Task, and Deadline Extraction

Description

Detect structured work signals from unstructured text and attachments.

Extractable entities

- dates,
- deadlines,
- hearing references,
- filing references,
- notice periods,
- tasks,
- assignees,
- document mentions,
- counterparties,
- legal trigger events.

Key value

This feature addresses a major legal risk area because deadline-related mistakes are repeatedly identified as a major source of malpractice exposure and operational failure.

11.5 Matter Graph / Legal Memory

Description

Store all extracted and linked information in a matter-centric data model.

Data entities

- Matter
- Client
- Contact / Party
- Channel / Thread
- Document

- Task
- Deadline
- Event
- Approval
- Knowledge Item
- Playbook Reference

This approach mirrors the value of legal-context-aware data layers promoted in legal AI systems and supports long-term retrieval and orchestration.

Why it matters

Instead of treating messages as isolated inputs, the system builds persistent legal context over time. That enables better triage, better Q&A, and more reliable timelines.

11.6 Timeline Builder

Description

Generate a source-linked chronology of matter events based on communications, documents, and extracted entities.

Output format

- date,
- event title,
- event summary,
- source references,
- linked documents,
- associated contacts/parties.

User value

Attorneys can rapidly understand the sequence of facts in a matter without manually re-reading all threads and attachments. This is especially valuable in litigation, disputes, and investigations.

11.7 Knowledge Retrieval / Grounded Legal Q&A

Description

Provide AI-assisted answers grounded in approved internal knowledge sources rather than generic, unsupported text generation.

Sources

- SharePoint folders,
- Google Drive folders,
- policy docs,
- legal playbooks,
- prior matter notes or approved templates,
- internal knowledge bases.

Requirements

- every answer must cite the internal sources used,
- access must respect repository permissions,
- responses should be explicit when the answer is incomplete or ambiguous.

This design is important given ABA guidance around competence, confidentiality, and supervision when lawyers use generative AI tools.

11.8 Action Layer (Approval-first)

Description

Allow users to turn extracted insights into actions.

v1 supported actions

- create task,
- assign owner,
- set due date,
- save matter note,

- create reminder,
- draft reply,
- flag urgent issue.

Control principle

No substantive outbound communication should be automatically sent in v1 without user review. This aligns with legal ethics expectations around human supervision and accountability for AI-generated work product

---

11.9 Auditability and Governance

Description

The system must provide legally appropriate visibility into how outputs were generated and what actions were taken.

Required controls

- audit log of ingestion, classification, extraction, and action events,
- visible source links behind summaries and extractions,
- user approvals for sensitive actions,
- role-based access control,
- data segregation by workspace/matter,
- configurable retention policies.

---

12) Functional Requirements

FR-1: Ingestion

The platform shall ingest messages and metadata from supported communication systems and documents from supported repositories.

FR-2: Normalization

The platform shall normalize inbound data into a common schema across email, chat, and document sources. This is necessary to support unified querying, extraction, and timeline generation across fragmented sources.

FR-3: Classification

The platform shall classify inbound legal work into predefined categories and assign urgency and routing suggestions.

FR-4: Summarization

The platform shall summarize multi-message threads and produce structured outputs including facts, issues, actions, and open questions.

FR-5: Extraction

The platform shall extract legal events, deadlines, tasks, assignees, named entities, and document references.

FR-6: Matter Linking

The platform shall associate communications and documents with an existing matter or propose a new matter grouping when none exists. This is fundamental to turning fragmented communication into structured legal work.

FR-7: Timeline Generation

The platform shall provide timeline generation for each matter with date-sorted events and traceable source references.

FR-8: Grounded Q&A

The platform shall answer questions over connected legal knowledge sources with citations to the underlying internal materials.

FR-9: Action Execution

The platform shall allow users to create tasks, reminders, and draft responses from extracted content.

FR-10: Approval Workflow

The platform shall require user approval for sensitive or externally visible actions.

FR-11: Auditability

The platform shall maintain a persistent audit trail of data ingestion, model outputs, approvals, and actions.

---

13) Non-Functional Requirements

13.1 Security

The platform must support secure authentication, role-aware authorization, and data isolation by tenant/workspace/matter. Legal users are highly sensitive to confidentiality and client data handling obligations.

13.2 Privacy

The platform must be explicit about whether data is processed by hosted LLMs, local models, or enterprise AI services, and should support privacy-preserving deployment options over time. Legal AI vendors emphasize privacy and enterprise-grade handling as a core trust requirement.

13.3 Explainability

Every major AI output must include source traceability and confidence/uncertainty handling. This helps satisfy legal-user expectations of supervision and reliability.

13.4 Performance

The system should return thread summaries quickly enough to support real-world inbox and chat workflows, and should generate timelines incrementally rather than requiring expensive full recomputation on every matter view. This requirement is product-driven to fit user workflows emphasized by embedded AI products.

13.5 Extensibility

The architecture should allow new connectors, new extractors, and new workflow actions to be added easily, consistent with the ecosystem ideas behind MCP and API-based integration tooling.

---

14) User Flows

Flow 1: Email/Chat Intake

1. A legal-related email or chat message arrives.
2. Connector ingests it through API/MCP layer.
3. System classifies the request by type/urgency.
4. System summarizes the thread and extracts tasks/deadlines/entities.
5. Matter is linked or proposed.

6. User reviews, approves, and assigns.

Flow 2: Timeline Generation

1. User opens a matter.
2. System aggregates related communications, documents, and extracted events.
3. Timeline is generated with dates and source links
4. User filters to only deadlines, client events, filings, or communications.
5. User exports or uses the timeline internally.

Flow 3: Grounded Legal Q&A

1. User asks a question such as: “What is our NDA playbook position on indemnity caps?”
2. System queries approved repositories and matter-linked knowledge.
3. System returns an answer with source citations and linked documents.
4. User optionally drafts a reply or creates a note.

---

15) Success Metrics

Product Metrics

- Number of messages/emails successfully ingested per workspace.
- Percentage of inbound items automatically classified into a valid legal category.
- Time saved per matter for summary/timeline creation.
- Average time from intake to assignment.
- Number of extracted action items and deadlines accepted by users.

These metrics tie directly to the operational inefficiencies legal teams face around intake, deadlines, and fragmented workflows.

Trust Metrics

- Percentage of summaries with source-linked support.
- User acceptance rate for extracted deadlines/tasks.
- Number of approval overrides or corrections.

- Number of hallucination/error reports.

These are important because legal AI adoption is constrained by trust, ethics, and supervision requirements.

Business / Open-source Metrics

- GitHub stars/forks/contributors.
- Number of connector adapters added by community.
- Number of example deployments.
- Demo-to-user conversion for pilots.

These align with the open, extensible architecture enabled by MCP and integration-driven product design

---

16) Risks and Mitigations

Risk 1: Hallucinated or misleading summaries

LLM outputs may overstate certainty or infer facts not supported by source material. This is especially risky in legal settings.

Mitigation:

- require source-linked outputs,
- separate extracted facts from inferred recommendations,
- highlight uncertainty,
- require human review for sensitive actions.

Risk 2: Confidential data exposure

Legal data is highly sensitive and may involve client confidences and privileged communications. ABA guidance makes confidentiality a central obligation in AI use.

Mitigation:

- enterprise-grade auth and authorization,
- clear data boundaries,
- optional private deployment path,

- repository-level permission checks,
- audit logs.

Risk 3: Integration complexity

Connecting to too many systems at once can slow delivery and increase maintenance burden. Both Wordsmith and other workflow tools demonstrate value through staged integration expansion.

Mitigation:

- prioritize Microsoft + Slack + Drive-style core connectors first,
- use MCP and connector abstractions,
- defer low-value or hard-to-govern integrations.

Risk 4: Deadline extraction errors

Extracting deadlines from text is valuable but risky if wrong. Legal deadline commentary shows how serious missed deadlines can be.

Mitigation:

- mark extracted deadlines as "suggested" until confirmed,
- provide source snippets,
- support manual confirmation before calendaring.

---

17) MVP Release Plan

Phase 1 — Foundation

- authentication and workspace model,
- connector framework,
- Outlook / Teams / Slack / SharePoint basic ingestion,
- normalized data schema,
- matter data model.

Phase 2 — Core Intelligence

- intake classification,

- summarization,
- task/event/deadline extraction,
- matter linking,
- timeline view.

Phase 3 — Grounded Assistance

- repository indexing,
- legal Q&A with citations,
- draft reply generation,
- approval-first action flow.

Phase 4 — Hardening

- audit log,
- governance controls,
- admin screens,
- retention controls,
- extensibility docs for contributors

---

18) High-Level Architecture

Architecture Summary

CaseFlow MCP will use a layered architecture:

Plain Text

Connectors Layer

├– Outlook / Microsoft 365

├– Teams

├– Slack

├– SharePoint / OneDrive

└ Google Drive / Gmail (later in v1.x)

Normalization Layer

├– unified message schema

├– unified document schema

└ entity mapping (matter/client/party/thread)

Intelligence Layer

├– intake classifier

├– summarizer

├– task/event/deadline extractor

├– matter linker

└ timeline builder

Knowledge Layer

├– repository index

├– permission-aware retrieval

└ grounded legal Q&A

Action Layer

├– create task

├– assign owner

├– save note

├– create reminder

└ draft reply (approval-first)

Governance Layer

├– RBAC

├– audit logs

├– data retention

└ source traceability

Show more lines

This architecture is consistent with modern integration-first legal AI products and with the MCP model of exposing data, tools, and workflows to AI systems

---

19) Open-Source Strategy

The open-source version should focus on the most reusable, developer-friendly parts of the platform:

- connector abstraction layer,
- legal matter/event schema,
- extraction pipelines,
- timeline engine,
- grounding framework for legal knowledge,
- lightweight web UI,
- sample mock datasets,
- contributor guide for adding new connectors and extractors.

This creates real open-source value because most public "legal AI" repos are document-chat demos rather than workflow-centric legal operating layers. The MCP-first architecture also makes extensibility a strong differentiator

---

---

21) Final Product Positioning

Positioning Statement

CaseFlow MCP is an open-source legal operations layer that connects Outlook, Teams, Slack, and legal knowledge repositories through MCP and API connectors to automatically triage requests, extract deadlines and legal events, build matter timelines, and give attorneys a single, grounded command center for legal work.

Why this positioning is strong

It is narrower and more achievable than a full Filevine replacement, but broader and more differentiated than a simple message summarizer. It directly addresses legal intake fragmentation, timeline reconstruction, deadline risk, and knowledge retrieval — all of which are validated pain areas in the current legal-tech market

---

If you want, I can do the **next best step** and turn this into one of these:

1. **Engineering architecture doc** (services, DB schema, MCP connectors, APIs)
2. **GitHub repo blueprint** (folders, modules, README, roadmap)
3. **Startup pitch deck content**
4. **Feature-by-feature MVP backlog in Jira style**

If you want the most practical next step, I'd suggest:

"Now build the system architecture + repo structure for this PRD."

How our platform is diff compared to existing software’s in the market

If you build it the way we defined, then your platform is different because:

- It is **matter-centric**, not tool-centric
- It converts communication into **structured legal work**
- It is a **platform layer**, not just an assistant
- It focuses heavily on **timeline + deadlines**, not just drafting

👉 The real difference is **what your system revolves around**

🧠 **Core Concept Difference**

Wordsmith

- Lives *inside tools* (Slack, Outlook, Word)
- Helps users in the moment
- Focus = **assist workflows**

👉 Think:

"AI helping me write or respond faster"

Your Platform

- Sits *above all tools*
- Pulls everything together into one system
- Focus = **organize legal work automatically**

👉 Think:

"System that organizes and tracks all legal work from communication"

🔥 **The Real Difference (in one line)**

Wordsmith

"Help me work better inside tools"

Your Platform

"Turn everything I receive into structured, trackable legal work"

✅ **4 Clear Differentiation Pillars**

1. Matter-Centric vs Tool-Centric

Wordsmith:

- Works inside Slack, Word, Outlook
- Each message or document is handled separately

Your Platform:

- Everything belongs to a **matter**
- Emails, chats, documents → all linked

👉 You are building:

A **case-level system**, not a message-level assistant

---

2. Communication → Work Engine (BIGGEST DIFFERENCE)

Wordsmith:

- Summarizes
- Drafts replies
- Answers questions

Your Platform:

- Converts messages into:
    - tasks
    - deadlines
    - events
    - structured actions

Example:

Wordsmith:

"Here’s a summary of the email"

Your Platform:

"This email creates:

- Task: review contract
- Deadline: June 30
- Matter: ACME dispute
- Event: breach notice received"

👉 This is a **huge shift**
You are creating **workflows automatically**, not just text outputs

---

3. Timeline + Deadline Focus

Wordsmith:

- Mostly focused on:
  - drafting
  - contracts
  - answering legal questions

Your Platform:

- Focuses on:
  - event reconstruction
  - timeline building
  - deadline detection
  - legal triggers

👉 Example output:

Plain Text

May 2 – Client complaint received

May 5 – Contract signed

May 10 – Breach reported

May 15 – Response due

Show more lines

👉 This solves:

- missed deadlines
- lack of visibility
- messy case history

This is where your product becomes **very powerful**

---

4. Platform vs Product

Wordsmith:

- Fixed SaaS product
- Predefined workflows
- Closed system

Your Platform:

- Open-source
- Modular
- Extensible

Meaning:

Others can add:

- new connectors
- new workflows
- new extractors

- new legal logic

👉 You are building:

A **legal infrastructure layer**

🎯 How you should POSITION it

❌ Weak positioning (too similar)

- AI for lawyers
- Legal assistant
- Slack/Outlook copilot

✅ Correct positioning

"A platform that converts fragmented legal communication into structured matters, timelines, deadlines, and actions."

✅ Best positioning (very strong)

"A legal work intelligence system that turns emails, chats, and documents into structured case data, timelines, and trackable legal workflows."

🧩 Mental Model

Existing tools

Plain Text

User → Tool (Slack/Outlook) → AI helps

` `

Show more lines

Your platform

Plain Text

All tools → Unified system → Structured matter → Timeline + Tasks

Show more lines

🚀 Final takeaway

You are NOT building:

- another chatbot
- another legal assistant
- another contract AI tool

You ARE building:

👉 A Legal Work Intelligence Engine

That:

- listens to everything
- understands it
- structures it
- tracks it
- and helps lawyers never lose context again

✅ Super Simple Pitch (use this anywhere)

“Instead of helping lawyers reply faster, we help them never lose track of what actually needs to be done.”